

*Cover design by Aurelia Bunesu*  
*Thanks to Fernando Batista and Matt Post for their helpful notes about creating a hand-*  
*book*  
*Handbook assembled by Swapna Somasundaran*



## Meal Info

---

The main conference registration includes

- Welcome Reception
- Two Poster Session dinners
- Two breaks each day with snacks

## WIFI access

---

**Username:** XXXXXXXXXXXXX

**Password:** XXXXXXXXXXXXX

## Links

---

Web site: <http://www.emnlp2016.net/>

Contact email: [contact@emnlp2016.net](mailto:contact@emnlp2016.net)

Find us on facebook: <https://www.facebook.com/EMNLP-2016-715509195252840/>

Follow us on twitter: <https://twitter.com/emnlp2016>

Conference4me app: XXXXXXXXXXXXX



## Contents

<b>1</b>	<b>Conference Information</b>	<b>1</b>
	Preface by the General Chair . . . . .	1
	Preface by the Program Committee Co-Chairs . . . . .	2
	Organizing Committee . . . . .	3
	Program Committee . . . . .	4
	Social events . . . . .	6
<b>2</b>	<b>Tutorials: Tuesday, November 1</b>	<b>7</b>
	<b>T1:</b> Practical Neural Networks for NLP: From Theory to Code . . . . .	8
	<b>T2:</b> Advanced Markov Logic Techniques for Scalable Joint Inference in NLP . . . . .	10
	<b>T3:</b> Lifelong Machine Learning for Natural Language Processing . . . . .	12
	<b>T4:</b> Neural Networks for Sentiment Analysis . . . . .	13
	<b>T5:</b> Continuous Vector Spaces for Cross-language NLP Applications . . . . .	15
	<b>T6:</b> Methods and Theories for Large-scale Structured Prediction . . . . .	16
<b>3</b>	<b>Workshops</b>	<b>19</b>
	<b>W1:</b> NLPSM . . . . .	20
	<b>W2:</b> SPNLP . . . . .	22
	<b>W3:</b> CSS . . . . .	24
	<b>W4:</b> LOUHI . . . . .	26
	<b>W5:</b> CNS . . . . .	28
	<b>W6:</b> UBLP . . . . .	29
	<b>Author index</b>	<b>31</b>



## Conference Information

### Preface by the General Chair

---

MOntH 25, 2016

XXXXXXXXXXXXXXXXXXXX XXXXXXXXXXXXXXX  
XXXXXXXXXXXXXXXXXXXX

<sup>1</sup>

*Jian Su*  
EMNLP 2016 General Chair

---

<sup>1</sup>Conference website: <http://www.emnlp2016.net>

## **Preface by the Program Committee Co-Chairs**

---

Month XX, 2016

XXX XXX XXX XXXX

EMNLP 2015 Program Committee Co-Chairs



## Organizing Committee

---

### General Chair

Jian Su, Institute for Infocomm Research - I2R

### Program co-Chairs

Xavier Carreras, Xerox Research Centre Europe  
Kevin Duh, Johns Hopkins University

### Workshops co-Chairs

Annie Louis, University of Essex  
Greg Kondrak, University of Alberta

### Tutorial co-Chairs

Bishan Yang, Carnegie Mellon University  
Rebecca Hwa, University of Pittsburgh

### Publication co-Chairs

Siddharth Patwardhan, IBM  
Daniele Pighin, Google Inc.

### Publicity Chair

Saif M. Mohammad, National Research Council Canada

### Sponsorship Chairs

Michel Galley, Microsoft Research  
Hang Li, Huawei Technologies (ISC Representative for EMNLP)

### Student Scholarship Chair and Student Volunteer Coordinator

Vincent Ng, University of Texas at Dallas

### Handbook Chair

Swapna Somasundaran, Educational Testing Service

### Website Chair

Jackie C.K. Cheung, McGill University

### SIGDAT Liaison

Chris Callison-Burch, University of Pennsylvania

### Local Arrangements Chair

Priscilla Rasmussen, ACL Business Manager

## Program Committee

---

### Program co-Chairs

Xavier Carreras, Xerox Research Centre Europe  
Kevin Duh, Johns Hopkins University

### Area Chairs

#### *Information Extraction, Information Retrieval, and Question Answering*

Nate Chambers, United States Naval Academy  
Ruihong Huang, Texas A&M University  
Min-Yen Kan, National University of Singapore  
Alan Ritter, The Ohio State University  
Scott Wen-Tau Yih, Microsoft Research

#### *Language and Vision*

Sanja Fidler, University of Toronto  
Julia Hockenmaier, University of Illinois at Urbana-Champaign

#### *Linguistic Theories and Psycholinguistics*

Suzanne Stevenson, University of Toronto

#### *Machine Learning*

Guillaume Bouchard, University College London  
Kyunghyun Cho, New York University  
Kuzman Ganchev, Google  
Ariadna Quattoni, Xerox Research Centre Europe  
Eric Ringger, Facebook

#### *Machine Translation and Multilinguality*

John DeNero, UC Berkeley  
Alex Fraser, Ludwig-Maximilians-Universitat Munchen  
Yang Liu, Tsinghua University  
Dekai Wu, Hong Kong University of Science and Technology

#### *Segmentation, Tagging, and Parsing*

Michael Collins, Columbia University  
Liang Huang, Oregon State University  
Daisuke Kawahara, Kyoto University  
Andre Martins, Priberam Labs

#### *Semantics*

Yoav Artzi, Cornell Tech  
Georgiana Dinu, IBM  
Ed Grefenstette, Google DeepMind  
Ray Mooney, University of Texas at Austin  
Laura Rimell, University of Cambridge

*Sentiment Analysis and Opinion Mining*

Dirk Hovy, University of Copenhagen  
Bing Liu, University of Illinois at Chicago  
Saif Mohammad, National Research Council Canada

*Social Media and Computational Social Science*

Tim Baldwin, University of Melbourne  
Smaranda Muresan, Columbia University

*Spoken Language Processing*

Brian Roark, Google  
Geoff Zweig, Microsoft Research

*Summarization, Generation, Discourse, Dialogue*

Manfred Stede, Potsdam University  
Michael Strube, Heidelberg Institute for Theoretical Studies  
Lucy Vanderwende, Microsoft Research  
Wei Xu, University of Pennsylvania

*Text Mining and NLP Applications*

Preslav Nakov, Qatar Computing Research Institute  
Vivi Nastase, FBK  
David Smith, Northeastern University  
Joel Tetreault, Yahoo! Labs

## **Social events**

---

XXXXXXXXXXXXXXXXX

### **Welcome Reception**

XXXX WHEN AND WHERE???

### **Farewell Drink**

Monday, September 21th at 18:30 at Culturgest

## Tutorials: Tuesday, November 1

### 9:00 – 12:30 **Morning Tutorials**

Practical Neural Networks for NLP: From Theory to Code

*Chris Dyer, Yoav Goldberg, and Graham Neubig*

*TBD TBD TBD TBD TBD TBD*

Advanced Markov Logic Techniques for Scalable Joint Inference in NLP

*Deepak Venugopal, Vibhav Gogate, and Vincent Ng*

*TBD*

Lifelong Machine Learning for Natural Language Processing

*Zhiyuan Chen and Bing Liu*

*TBD TBD TBD TBD TBD TBD*

### 12:30 – 14:00 **Lunch break**

### 14:00 – 17:30 **Afternoon Tutorials**

Neural Networks for Sentiment Analysis

*Yue Zhang and Duy Tin Vo*

*TBD*

Continuous Vector Spaces for Cross-language NLP Applications

*Rafael E. Banchs*

*TBD TBD TBD TBD TBD TBD*

Methods and Theories for Large-scale Structured Prediction

*Xu Sun and Yansong Feng*

*TBD*

# Practical Neural Networks for NLP: From Theory to Code

*Chris Dyer, Yoav Goldberg, and Graham Neubig*

Tuesday, November 1 – 09:00–12:30

TBD TBD TBD TBD TBD TBD

This tutorial aims to bring NLP researchers up to speed with the current techniques in deep learning and neural networks, and show them how they can turn their ideas into practical implementations. We will start with simple classification models (logistic regression and multilayer perceptrons) and cover more advanced patterns that come up in NLP such as recurrent networks for sequence tagging and prediction problems, structured networks (e.g., compositional architectures based on syntax trees), structured output spaces (sequences and trees), attention for sequence-to-sequence transduction, and feature induction for complex algorithm states. A particular emphasis will be on learning to represent complex objects as recursive compositions of simpler objects. This representation will reflect characterize standard objects in NLP, such as the composition of characters and morphemes into words, and words into sentences and documents. In addition, new opportunities such as learning to embed “algorithm states” such as those used in transition-based parsing and other sequential structured prediction models (for which effective features may be difficult to engineer by hand) will be covered.

Everything in the tutorial will be grounded in code — we will show how to program seemingly complex neural-net models using toolkits based on the computation-graph formalism. Computation graphs decompose complex computations into a DAG, with nodes representing inputs, target outputs, parameters, or (sub)differentiable functions (e.g., “tanh”, “matrix multiply”, and “softmax”), and edges represent data dependencies. These graphs can be run “forward” to make predictions and compute errors (e.g., log loss, squared error) and then

---

**Chris Dyer (CMU)** Chris’s research looks at the intersection between machine learning and natural language processing, in particular focusing on representation learning, structured models, and multilinguality. He has received best paper nominations and awards at NAACL, EMNLP, and ACL. He is the author of the CNN toolkit, which was used by the Jelinek 2016 Summer Workshop on neural machine translation.

**Yoav Goldberg (Bar Ilan)** Yoav’s research focuses around structure prediction problems in natural language, in particular prediction of syntactic representations, greedy methods, semi-supervised learning and robust cross-domain performance. Recently, Yoav became involved in the application of neural-network based models to NLP problems, with two core focuses: better understanding the neural network building blocks; and the use of neural networks for structured problems. Yoav wrote a primer on neural-networks method in NLP, and is responsible for the python bindings of the CNN library.

**Graham Neubig (NAIST)** Graham’s research focuses on machine learning methods for language and speech processing. His work puts a particular focus on machine translation and speech translation, but also covers other NLP tasks such as tagging and syntactic/semantic parsing, and methods for learning from unlabeled data such as unsupervised, semi-supervised, or active learning. He is interested in neural network models that incorporate our linguistic intuitions, and is one of the core developers of the CNN toolkit.

"backward" to compute derivatives with respect to model parameters. In particular we'll cover the Python bindings of the CNN library. CNN has been designed from the ground up for NLP applications, dynamically structured NNs, rapid prototyping, and a transparent data and execution model.

# Advanced Markov Logic Techniques for Scalable Joint Inference in NLP

*Deepak Venugopal, Vibhav Gogate, and Vincent Ng*

Tuesday, November 1 – 09:00–12:30

TBD

In the early days of the statistical NLP era, many language processing tasks were tackled using the so-called pipeline architecture: the given task is broken into a series of sub-tasks such that the output of one sub-task is an input to the next sub-task in the sequence. The pipeline architecture is appealing for various reasons, including modularity, modeling convenience, and manageable computational complexity. However, it suffers from the error propagation problem: errors made in one sub-task are propagated to the next sub-task in the sequence, leading to poor accuracy on that sub-task, which in turn leads to more errors downstream. Another disadvantage associated with it is lack of feedback: errors made in a sub-task are often not corrected using knowledge uncovered while solving another sub-task down the pipeline.

Realizing these weaknesses, researchers have turned to joint inference approaches in recent years. One such approach involves the use of Markov logic, which is defined as a set of weighted first-order logic formulas and, at a high level, unifies first-order logic with probabilistic graphical models. It is an ideal modeling language (knowledge representation) for compactly representing relational and uncertain knowledge in NLP. In a typical use case of MLNs in NLP, the application designer describes the background knowledge using a few first-order logic sentences and then uses software packages such as *Alchemy*, *Tuffy*, and *Markov the beast* to perform learning and inference (prediction) over the MLN. However,

---

**Deepak Venugopal** completed his Ph.D. at the University of Texas at Dallas in 2015 after which he joined the department of computer science at the University of Memphis where he is currently an assistant professor. His main research interests are in Probabilistic Graphical Models, Statistical Relational Models and their applications in natural language processing and cyber-security. His research work has resulted in several key techniques for scalable inference and learning in Statistical Relational Models.

**Vibhav Gogate** is an Assistant Professor in the Computer Science Department at the University of Texas at Dallas. He got his Ph.D. at University of California, Irvine in 2009 and then did a two-year post-doc at University of Washington. His current research focus is on probabilistic graphical models and its first-order logic based extensions such as Markov logic. He is best known for his work on inference algorithms that combine the power of logic and probability including lifted probabilistic inference algorithms. He is the co-winner of the last two probabilistic inference competitions the 2010 UAI approximate inference challenge and the 2012 PASCAL probabilistic inference competition.

**Vincent Ng** is an Associate Professor of Computer Science and a member of the Human Language Technology Research Institute at the University of Texas at Dallas. He is best known for his work on coreference resolution in the NLP community. He has recently been collaborating with the co-authors of this proposal on modeling complex NLP tasks using Markov Logic.



despite its obvious advantages, over the years, researchers and practitioners have found it difficult to use MLNs effectively in many NLP applications. The main reason for this is that it is hard to scale inference and learning algorithms for MLNs to large datasets and complex models, that are typical in NLP.

In this tutorial, we will introduce the audience to recent advances in scaling up inference and learning in MLNs as well as new approaches to make MLNs a “black-box” for NLP applications (with only minor tuning required on the part of the user). Specifically, we will introduce attendees to a key idea that has emerged in the MLN research community over the last few years, lifted inference, which refers to inference techniques that take advantage of symmetries (e.g., synonyms), both exact and approximate, in the MLN. We will describe how these next-generation inference techniques can be used to perform effective joint inference. We will also present our new software package for inference and learning in MLNs, Alchemy 2.0, which is based on lifted inference, focusing primarily on how it can be used to scale up inference and learning in large models and datasets for applications such as semantic similarity determination, information extraction and question answering.

# Lifelong Machine Learning for Natural Language Processing

*Zhiyuan Chen and Bing Liu*

Tuesday, November 1 – 09:00–12:30

TBD TBD TBD TBD TBD TBD

Machine learning (ML) has been successfully used as a prevalent approach to solving numerous NLP problems. However, the classic ML paradigm learns in isolation. That is, given a dataset, an ML algorithm is executed on the dataset to produce a model without using any related or prior knowledge. Although this type of isolated learning is very useful, it also has serious limitations as it does not accumulate knowledge learned in the past and use the knowledge to help future learning, which is the hallmark of human learning and human intelligence. Lifelong machine learning (LML) aims to achieve this capability. Specifically, it aims to design and develop computational learning systems and algorithms that learn as humans do, i.e., retaining the results learned in the past, abstracting knowledge from them, and using the knowledge to help future learning. In this tutorial, we will introduce the existing research of LML and to show that LML is very suitable for NLP tasks and has potential to help NLP make major progresses.

---

**Zhiyuan Chen, Google** Zhiyuan Chen completed his Ph.D. at the University of Illinois at Chicago (UIC) and joined Google in 2016. His research interests include Text Mining, Machine Learning, and Statistical Natural Language Processing. He has designed several life-long learning algorithms to automatically mine valuable information from text documents. He has published more than 15 full research papers at premier conferences.

**Bing Liu, University of Illinois at Chicago (UIC)** Bing Liu is a professor of Computer Science at the University of Illinois at Chicago (UIC). He received his PhD in Artificial Intelligence from the University of Edinburgh. His current research interests include sentiment analysis, NLP, data mining, and machine learning. He has published extensively in top conferences and journals, and also three books. Two of his papers received test-of-time awards from KDD. He is a Fellow of ACM, AAAI, and IEEE.

## Neural Networks for Sentiment Analysis

Yue Zhang and Duy Tin Vo

Tuesday, November 1 – 14:00–17:30

TBD

Sentiment analysis has been a major research topic in natural language processing (NLP). Traditionally, the problem has been attacked using discrete models and manually-defined sparse features. Over the past few years, neural network models have received increased research efforts in most sub areas of sentiment analysis, giving highly promising results. A main reason is the capability of neural models to automatically learn dense features that capture subtle semantic information over words, sentences and documents, which are difficult to model using traditional discrete features based on words and ngram patterns. This tutorial gives an introduction to neural network models for sentiment analysis, discussing the mathematics of word embeddings, sequence models and tree structured models and their use in sentiment analysis on the word, sentence and document levels, and fine-grained sentiment analysis. The tutorial covers a range of neural network models (e.g. CNN, RNN, RecNN, LSTM) and their extensions, which are employed in four main subtasks of sentiment analysis:

- Sentiment-oriented embeddings;
- Sentence-level sentiment;
- Document-level sentiment;
- Fine-grained sentiment.

The content of the tutorial is divided into 3 sections of 1 hour each. We assume that the audience is familiar with linear algebra and basic neural network structures, introduce the mathematical details of the most typical models. First, we will introduce the sentiment

---

**Yue Zhang, Assistant Professor, Singapore University of Technology and Design** Yue Zhang is an assistant professor at Singapore University of Technology and Design (SUTD). Before joining SUTD, he worked as a postdoctoral research fellow at University of Cambridge. He received his DPhil and MSc degrees from University of Oxford, and BEng from Tsinghua University. His main research interest includes machine learning, syntactic parsing and text mining. Over the recent years, he worked intensively on deep learning and sentiment analysis. He serves in the standing review committee of TACL, and was area chair of COLING'14, NAACL'15 and EMNLP'15. Yue Zhang has given a tutorial at NAACL'10 and a tutorial at ACL'14.

**Duy Tin Vo, PhD Candidate, Singapore University of Technology and Design** Duy Tin Vo is a PhD candidate at Singapore University of Technology and Design (SUTD) under the supervision of Yue Zhang. He received his BEng from Cantho University. His current research focuses on applying machine learning and deep learning in sentiment analysis.

analysis task, basic concepts related to neural network models for sentiment analysis, and show detail approaches to integrate sentiment information into embeddings. Sentence-level models will be described in the second section. Finally, we will discuss neural network models use for document-level and fine-grained sentiment.

## Continuous Vector Spaces for Cross-language NLP Applications

*Rafael E. Banchs*

Tuesday, November 1 – 14:00–17:30

TBD TBD TBD TBD TBD TBD

The mathematical metaphor offered by the geometric concept of distance in vector spaces with respect to semantics and meaning has been proven to be useful in many monolingual natural language processing applications. There is also some recent and strong evidence that this paradigm can also be useful in the cross-language setting. In this tutorial, we present and discuss some of the most recent advances on exploiting the vector space model paradigm in specific cross-language natural language processing applications, along with a comprehensive review of the theoretical background behind them.

First, the tutorial introduces some fundamental concepts of distributional semantics and vector space models. More specifically, the concepts of distributional hypothesis and term-document matrices are revised, followed by a brief discussion on linear and non-linear dimensionality reduction techniques and their implications to the parallel distributed approach to semantic cognition. Next, some classical examples of using vector space models in monolingual natural language processing applications are presented. Specific examples in the areas of information retrieval, related term identification and semantic compositionality are described.

Then, the tutorial focuses its attention on the use of the vector space model paradigm in cross-language applications. To this end, some recent examples are presented and discussed in detail, addressing the specific problems of cross-language information retrieval, cross-language sentence matching, and machine translation. Some of the most recent developments in the area of Neural Machine Translation are also discussed.

Finally, the tutorial concludes with a discussion about current and future research problems related to the use of vector space models in cross-language settings. Future avenues for scientific research are described, with major emphasis on the extension from vector and matrix representations to tensors, as well as the problem of encoding word position information into the vector-based representations.

---

**Rafael E Banchs, Researcher, Institute for Infocomm Research, Singapore** Rafael is a Research Scientist at the Institute for Infocomm Research, in Singapore, where he leads the Dialogue Technology Lab of the Human Language Technology Department. His main area of research is focused on the construction and use of semantic representations to support different natural language processing applications, including machine translation, information retrieval, natural language understanding and chat-oriented dialogue.

## Methods and Theories for Large-scale Structured Prediction

*Xu Sun and Yansong Feng*

Tuesday, November 1 – 14:00-17:30

TBD

Many important NLP tasks are casted as structured prediction problems, and try to predict certain forms of structured output from the input. Examples of structured prediction include POS tagging, named entity recognition, PCFG parsing, dependency parsing, machine translation, and many others. When apply structured prediction to a specific NLP task, there are the following challenges:

1. Model selection: Among various models/algorithms with different characteristics, which one should we choose for a specific NLP task?
2. Training: How to train the model parameters effectively and efficiently?
3. Overfitting: To achieve good accuracy on test data, it is important to control the overfitting from the training data. How to control the overfitting risk for structured prediction?

This tutorial will provide a clear overview of recent advances in structured prediction methods and theories, and address the above issues when we apply structured prediction to NLP tasks. We will introduce large margin methods (e.g., perceptrons, MIRA), graphical models (e.g., CRFs), and deep learning methods (e.g., RNN, LSTM), and show the respective advantages and disadvantages for NLP applications. For the training algorithms, we will introduce online/ stochastic training methods, and we will introduce parallel online/stochastic

---

**Xu Sun, Associate Professor, Peking University** Xu Sun is currently an Associate Professor at the School of EECS, Peking University. He got PhD from The University of Tokyo in 2010. He worked at The University of Tokyo, Cornell University, and The Hong Kong Polytechnic University as research fellow/associate from 2010 to 2012. His research interests focus on natural language processing and machine learning, especially on structured natural language processing and structured prediction, and with related publications on ACL, NIPS, IJCAI, Comput. Linguist., IEEE TKDE, etc. He has been area chair of EMNLP; program committee member of ACL, IJCAI, AAAI, NAACL; journal reviewer of IEEE TPAMI, Comput. Linguist., IEEE T-ASL, and so on.

**Yansong Feng, Lecturer, Peking University** Yansong Feng is a Lecturer at the Institute of Computer Science and Technology, Peking University. He received his Ph.D. degree from ILCC, the University of Edinburgh and worked as a postdoctoral research assistant there in 2011. His research interests focus on applying machine learning techniques to natural language processing and information retrieval applications, such as information extraction, question answering, knowledge acquisition, and so on. His works have been published in top refereed international journals and conferences, such as IEEE TPAMI, ACL, VLDB, EMNLP, NAACL, AAAI, etc. He has served as program committee member of AAAI, EMNLP, NAACL, WWW; journal reviewer of IEEE TASLP, IEEE TKDE, TACL, and so on.

learning algorithms and theories to speed up the training (e.g., the Hogwild algorithm). For controlling the overfitting from training data, we will introduce the weight regularization methods, structure regularization, and implicit regularization methods.







# Workshop 1: NLPSM

## 4th International Workshop on Natural Language Processing for Social Media

---

Organizers: *Jane Yung-jen Hsu, Lun-Wei Ku, and Cheng-Te Li*

Venue: Salon J

**Tuesday, November 1, 2016**

08:50–09:00 **Opening**

09:00–10:00 **Keynote Speech: Cristian Danescu-Niculescu-Mizil**

10:00–10:30 **Short Paper Session I**

10:00–10:15 Identifying and Categorizing Disaster-Related Tweets  
*Kevin Stowe, Michael J. Paul, Martha Palmer, Leysia Palen, and Kenneth Anderson*

10:15–10:30 Identifying Eyewitness News-worthy Events on Twitter  
*Erika Doggett and Alejandro Cantarero*

10:30–11:00 **Coffee Break**

11:00–12:15 **Short Paper Session II**

11:00–11:15 Why Do They Leave: Modeling Participation in Online Depression Forums  
*Farig Sadeque, Ted Pedersen, Thamar Solorio, Prasha Shrestha, Nicolas Rey-Villamizar, and Steven Bethard*

11:15–11:30 Twitter at the Grammys: A Social Media Corpus for Entity Linking and Disambiguation  
*Mark Dredze, Nicholas Andrews, and Jay DeYoung*

11:30–11:45 Steps Toward Automatic Understanding of the Function of Affective Language in Support Groups  
*Amit Navindgi, Caroline Brun, Cécile Boulard Masson, and Scott Nowson*

11:45–12:00 Detecting Social Roles in Twitter  
*Sunghwan Mac Kim, Stephen Wan, and Cecile Paris*

12:00–12:15 Identifying Sensible Participants in Online Discussions  
*Siddharth Jain*

12:30–14:00 **Lunch**

14:00–15:30 **Long Paper Session I**

14:00–14:30 emoji2vec: Learning Emoji Representations from their Description  
*Ben Eisner, Tim Rocktäschel, Isabelle Augenstein, Matko Bosnjak, and Sebastian Riedel*

14:30–15:00 Learning Latent Local Conversation Modes for Predicting Comment Endorsement in Online Discussions  
*Hao Fang, Hao Cheng, and Mari Ostendorf*

- 15:00–15:30 Witness Identification in Twitter  
*Rui Fang, Armineh Nourbakhsh, XIAOMO LIU, Sameena Shah, and Quanzhi Li*
- 15:30–16:00 **Coffee Break / Poster Session and Discussion**
- 16:00–17:30 **Long Paper Session II**
- 16:00–16:30 How Do I Look? Publicity Mining From Distributed Keyword Representation of Socially Infused News Articles  
*Yu-Lun Hsieh, Yung-Chun Chang, Chun-Han Chu, and Wen-Lian Hsu*
- 16:30–17:00 Hierarchical Character-Word Models for Language Identification  
*Aaron Jaech, George Mulcaire, Shobhit Hathi, Mari Ostendorf, and Noah A. Smith*
- 17:00–17:30 Human versus Machine Attention in Document Classification: A Dataset with Crowdsourced Annotations  
*Nikolaos Pappas and Andrei Popescu-Belis*
- 17:30–17:40 **Award and Closing**

## Workshop 2: SPNLP

### Structured Prediction for NLP

---

Organizers: *Kai-Wei Chang, Alexander Rush, Vivek Srikumar, and Ming-Wei Chang*

Venue: The Gallery

### Saturday, November 5, 2016

#### 09:00–10:30 **Section 1**

09:00–09:15 *Welcome (Organizers)*

09:15–10:00 *Invited Talk (Kristina Toutanova)*

10:00–10:30 *Inside-Outside and Forward-Backward Algorithms Are Just Backprop  
(tutorial paper)  
Jason Eisner*

#### 10:30–11:00 **Coffee break**

#### 11:00–12:30 **Section 2**

11:00–11:45 *Invited Talk (Andrew McCallum)*

11:45–12:30 *Panel (Panel: Moderator - Hal Daume)*

#### 12:30–14:00 **Lunch**

#### 14:00–15:30 **Section 3**

- Research on attention memory networks as a model for learning natural language inference  
*Degen Huang*
- A Joint Model of Rhetorical Discourse Structure and Summarization  
*Naman Goyal and Jacob Eisenstein*
- Posterior regularization for Joint Modeling of Multiple Structured Prediction Tasks with Soft Constraints  
*Kartik Goyal and Chris Dyer*
- A Study of Imitation Learning Methods for Semantic Role Labeling  
*Travis Wolfe, Mark Dredze, and Benjamin Van Durme*
- Introducing DRAIL – a Step Towards Declarative Deep Relational Learning  
*Xiao Zhang, María Leonor Pacheco, Chang Li, and Dan Goldwasser*

14:45–15:30 *Invited Talk (Raquel Urtasun)*

#### 14:00–14:45 **Poster**

#### 15:30–16:00 **Coffee Break**

#### 16:00–17:50 **Section 4**

16:00–16:45 *Invited Talk (Dzmitry Bahdanau)*

16:45–17:15 Unsupervised Neural Hidden Markov Models

*Ke M. Tran, Yonatan Bisk, Ashish Vaswani, Daniel Marcu, and Kevin Knight*

17:15–17:30 **Closing**

## Workshop 3: CSS

### Natural Language Processing and Computational Social Science

---

Organizers: *Dirk Hovy, David Bamman, A. Seza Dogruoz, Jacob Eisenstein, David Jurgens, Brendan O'Connor, Alice Oh, Oren Tsur, and Svitlana Volkova*

Venue: Broken Spoke

**Saturday, November 5, 2016**

09:00–10:30 **Session 1**

09:00–09:15 *Welcome (Workshop organizers)*

09:15–10:00 *Invited talk (James Pennebaker)*

10:00–10:30 **Doctoral consortium presentations**

10:30–11:00 **coffee break**

11:00–12:30 **Session 2**

11:00–11:45 *Invited talk (Hanna Wallach)*

11:45–12:30 *Invited talk (Cristian Danescu-Niculescu-Mizil)*

12:30–14:00 **Lunch break**

14:00–15:30 **Session 3**

14:00–14:45 *Invited talk (Jason Baldridge)*

14:45–15:30 **1-minute poster madness**

- Relating semantic similarity and semantic association to how humans label other people  
*Kenneth Joseph and Kathleen M. Carley*
- Identifying News from Tweets  
*Jesse Freitas and Heng Ji*
- Obfuscating Gender in Social Media Writing  
*Shravana Reddy and Kevin Knight*
- Social Proof: The Impact of Author Traits on Influence Detection  
*Sara Rosenthal and Kathy McKeown*
- Generating Politically-Relevant Event Data  
*John Beielor*
- User profiling with geo-located posts and demographic data  
*Adam Poulston, Mark Stevenson, and Kalina Bontcheva*
- Gov2Vec: Learning Distributed Representations of Institutions and Their Legal Text  
*John J. Nay*

- #WhoAmI in 160 Characters? Classifying Social Identities Based on Twitter Profile Descriptions  
*Anna Priante, Djoerd Hiemstra, Tijs van den Broek, Aaqib Saeed, Michel Ehrenhard, and Ariana Need*
- Identifying Stance by Analyzing Political Discourse on Twitter  
*Kristen Johnson and Dan Goldwasser*
- Learning Linguistic Descriptors of User Roles in Online Communities  
*Alex Wang, William L. Hamilton, and Jure Leskovec*
- The Effects of Data Collection Methods in Twitter  
*Sunghwan Mac Kim, Stephen Wan, Cecile Paris, Jin Brian, and Bella Robinson*
- Expressions of Anxiety in Political Texts  
*Ludovic Rheault*
- Constructing an Annotated Corpus for Protest Event Mining  
*Peter Makarov, Jasmine Lorenzini, and Hanspeter Kriesi*
- Demographer: Extremely Simple Name Demographics  
*Rebecca Knowles, Josh Carroll, and Mark Dredze*
- Bag of What? Simple Noun Phrase Extraction for Text Analysis  
*Abram Handler, Matthew Denny, Hanna Wallach, and Brendan O'Connor*
- News Sentiment and Cross-Country Fluctuations  
*Samuel Fraiberger*
- The Clinical Panel: Leveraging Psychological Expertise During NLP Research  
*Glen Coppersmith, Kristy Hollingshead, H. Andrew Schwartz, Molly Ireland, Rebecca Resnik, Kate Loveys, April Foreman, and Loring Ingraham*
- Are You a Racist or Am I Seeing Things? Annotator Influence on Hate Speech Detection on Twitter  
*Zeerak Waseem*
- Disentangling Topic Models: A Cross-cultural Analysis of Personal Values through Words  
*Steven Wilson, Rada Mihalcea, Ryan Boyd, and James Pennebaker*

15:30–16:00 **coffee break**

16:00–16:45 **posters**

16:00–17:30 **Session 4**

16:45–17:30 *Invited talk (Molly Roberts)*

## Workshop 4: LOUHI

### Louhi 2016: the Seventh International Workshop on Health Text Mining and Information Analysis

---

Organizers: Cyril Grouin, Thierry Hamon, Aurélie Névéol, and Pierre Zweigenbaum

Venue: TBD

**Saturday, November 5, 2016**

**09:00–10:15 Session I - Machine-Learning**

- 09:00–09:25 An Investigation of Recurrent Neural Architectures for Drug Name Recognition  
*Raghavendra Chalapathy, Ehsan Zare Borzeshi, and Massimo Piccardi*
- 09:25–09:50 Clinical Text Prediction with Numerically Grounded Conditional Language Models  
*Georgios Spithourakis, Steffen Petersen, and Sebastian Riedel*
- 09:50–10:15 Modelling Radiological Language with Bidirectional Long Short-Term Memory Networks  
*Savelie Cornegruta, Robert Bakewell, Samuel Withey, and Giovanni Montana*

**10:15–10:30 Session II - Boosters**

**11:00–12:30 Session III - Posters**

- Data Resource Acquisition from People at Various Stages of Cognitive Decline — Design and Exploration Considerations  
*Dimitrios Kokkinakis, Kristina Lundholm Fors, and Arto Nordlund*
- Analysis of Anxious Word Usage on Online Health Forums  
*Nicolas Rey-Villamizar, Prasha Shrestha, Farig Sadeque, Steven Bethard, Ted Pedersen, Arjun Mukherjee, and Tamar Solorio*
- Retrofitting Word Vectors of MeSH Terms to Improve Semantic Similarity Measures  
*Zhiguo Yu, Trevor Cohen, Byron Wallace, Elmer Bernstam, and Todd Johnson*
- Unsupervised Resolution of Acronyms and Abbreviations in Nursing Notes Using Document-Level Context Models  
*Katrin Kirchhoff and Anne M. Turner*
- Low-resource OCR error detection and correction in French Clinical Texts  
*Eva D'hondt, Cyril Grouin, and Brigitte Grau*
- Citation Analysis with Neural Attention Models  
*Tsendsuren Munkhdalai, John Lalor, and Hong Yu*
- Replicability of Research in Biomedical Natural Language Processing: a pilot evaluation for a coding task  
*Aurelie Neveol, Kevin Cohen, Cyril Grouin, and Aude Robert*



12:30–14:00 **Lunch break**

14:00–15:30 **Session IV - Invited talk**

14:00–15:30 NLP and Online Health Reports: What do we say and what do we mean?  
*Nigel Collier*

16:00–17:15 **Session V - NLP for literature and clinical documents**

16:00–16:25 Leveraging coreference to identify arms in medical abstracts: An experimental study  
*Elisa Ferracane, Iain Marshall, Byron C. Wallace, and Katrin Erk*

16:25–16:50 Hybrid methods for ICD-10 coding of death certificates  
*Pierre Zweigenbaum and Thomas Lavergne*

16:50–17:15 Exploring Query Expansion for Entity Searches in PubMed  
*Chung-Chi Huang and Zhiyong Lu*

## Workshop 5: CNS Computing News Storylines

---

Organizers: *Tommaso Caselli, Ben Miller, Marieke van Erp, Piek Vossen, and David Caswell*

Venue: Moody

### Friday, November 5, 2016

09:00–10:30 **Session 1:**

09:10–10:10 *Bridging the Gap between Event Macro-structures and Event Micro-structures*  
*(Ed Hovy, Language Technologies Institute, CMU)*

09:00–09:10 **Welcome and Opening Remarks**

10:30–11:00 **Coffee Break**

11:00–12:30 **Session 2:**

11:00–11:25 *Computable News Ecosystems: Roles for Humans and Machines*  
*David Caswell*

11:25–11:50 *Storyline detection and tracking using Dynamic Latent Dirichlet Allocation*  
*Daniel Bruggemann, Yannik Hermey, Carsten Orth, Darius Schneider, Stefan Selzer, and Gerasimos Spanakis*

11:50–12:15 *Real-time News Story Detection and Tracking with Hashtags*  
*Gevorg Poghosyan and Georgiana Ifrim*

12:15–12:30 *Nonparametric Bayesian Storyline Detection from Microtexts*  
*Vinodh Krishnan and Jacob Eisenstein*

12:30–14:00 **Lunch**

14:00–15:30 **Session 3:**

14:00–14:25 *Automatic Identification of Narrative Diegesis and Point of View*  
*Joshua Eisenberg and Mark Finlayson*

14:25–14:50 *Richer Event Description: Integrating event coreference with temporal, causal and bridging annotation*  
*Tim O’Gorman, Kristin Wright-Bettner, and Martha Palmer*

14:50–15:15 *NASTEa: Investigating Narrative Schemas through Annotated Entities*  
*Dan Simonson and Anthony Davis*

15:15–15:30 *The Storyline Annotation and Representation Scheme (StaR): A Proposal*  
*Tommaso Caselli and Piek Vossen*

16:00–17:00 **Discussion and Closing Remarks**

16:00–17:00 **Session 4:**

## Workshop 6: UBLP

### Uphill Battles in Language Processing: Scaling early achievements to robust methods

---

Organizers: Annie Louis, Michael Roth, Bonnie Webber, Michael White, and  
Luke Zettlemoyer

Venue: Continental Ballroom II

### Saturday, November 5, 2016

09:00–10:30 **Text Understanding**

09:00–10:20 *Invited talks, followed by discussion (Hal Daume III, Andrew Kehler, Chris Manning, Marie-Catherine de Marneffe)*

10:20–10:30 **Poster Boasters**

- An Analysis of Prerequisite Skills for Reading Comprehension  
*Saku Sugawara and Akiko Aizawa*
- Bridging the gap between computable and expressive event representations in Social Media  
*Darina Benikova and Torsten Zesch*
- Statistical Script Learning with Recurrent Neural Networks  
*Karl Pichotta and Raymond Mooney*
- Moving away from semantic overfitting in disambiguation datasets  
*Marten Postma, Filip Ilievski, Piek Vossen, and Marieke van Erp*
- Unsupervised Event Coreference for Abstract Words  
*Dheeraj Rajagopal, Eduard Hovy, and Teruko Mitamura*
- Towards Broad-coverage Meaning Representation: The Case of Comparison Structures  
*Omid Bakhshandeh and James Allen*

10:30–11:00 **Coffee break**

11:00–12:30 **Generation**

11:00–12:20 *Invited talks, followed by discussion (David DeVault, Mark Liberman, Diane Litman, Amanda Stent)*

12:20–12:30 **Poster Boasters 2**

- DialPort: A General Framework for Aggregating Dialog Systems  
*Tiancheng Zhao, Kyusong Lee, and Maxine Eskenazi*
- C2D2E2: Using Call Centers to Motivate the Use of Dialog and Diarization in Entity Extraction  
*Ken Church, Weizhong Zhu, and Jason Pelecanos*

- Visualizing the Content of a Children’s Story in a Virtual World: Lessons Learned  
*Quynh Ngoc Thi Do, Steven Bethard, and Marie-Francine Moens*
- Stylistic Transfer in Natural Language Generation Systems Using Recurrent Neural Networks  
*Jad Kabbara and Jackie Chi Kit Cheung*
- Using Language Groundings for Context-Sensitive Text Prediction  
*Timothy Lewis, Cynthia Matuszek, Amy Hurst, and Matthew Taylor*
- Towards a continuous modeling of natural language domains  
*Sebastian Ruder, Parsa Ghaffari, and John G. Breslin*

12:30–14:00 **Lunch break**

14:00–15:20 **Dialogue and Speech**

14:00–15:20 *Invited talks, followed by discussion (Ioannis Konstas, Kathleen McKeown, Margaret Mitchell, Donia Scott)*

15:20–16:00 **Coffee break + poster session**

16:00–16:30 **Poster session (continued)**

16:30–17:50 **Grounded Language**

16:30–17:50 *Invited talks, followed by discussion (James Allen, Joyce Chai, Yejin Choi, Mark Steedman)*

## Index

- Anderson, Kenneth, 20  
Andrews, Nicholas, 20  
Artzi, Yoav, 4  
Augenstein, Isabelle, 20  
  
Bakewell, Robert, 26  
Baldwin, Tim, 5  
Bamman, David, 24  
Banchs, Rafael E., 7, 15  
Beielser, John, 24  
Bernstam, Elmer, 26  
Bethard, Steven, 20, 26  
Bisk, Yonatan, 23  
Bontcheva, Kalina, 24  
Bosnjak, Matko, 20  
Bouchard, Guillaume, 4  
Boulard Masson, Cécile, 20  
Boyd, Ryan, 25  
Brian, Jin, 25  
Broek, Tijs van den, 25  
Bruggermann, Daniel, 28  
Brun, Caroline, 20  
  
Callison-Burch, Chris, 2, 3  
Cantarero, Alejandro, 20  
Carley, Kathleen M., 24  
Carreras, Xavier, 3, 4  
Carroll, Josh, 25  
Caselli, Tommaso, 28  
Caswell, David, 28  
Chalapathy, Raghavendra, 26  
Chambers, Nate, 4  
Chang, Kai-Wei, 22  
Chang, Ming-Wei, 22  
  
Chang, Yung-Chun, 21  
Chen, Zhiyuan, 7, 12  
Cheng, Hao, 20  
Cheung, Jackie C.K., 3  
Cho, Kyunghyun, 4  
Chu, Chun-Han, 21  
Cohen, Kevin, 26  
Cohen, Trevor, 26  
Collier, Nigel, 27  
Collins, Michael, 4  
Coppersmith, Glen, 25  
Cornegruta, Savelie, 26  
  
D'hondt, Eva, 26  
Davis, Anthony, 28  
DeNero, John, 4  
Denny, Matthew, 25  
DeYoung, Jay, 20  
Dinu, Georgiana, 4  
Doggett, Erika, 20  
Dogruoz, A. Seza, 24  
Dredze, Mark, 20, 22, 25  
Duh, Kevin, 3, 4  
Dyer, Chris, 7, 8, 22  
  
E Banchs, Rafael, 15  
Ehrenhard, Michel, 25  
Eisenberg, Joshua, 28  
Eisenstein, Jacob, 22, 24, 28  
Eisner, Ben, 20  
Eisner, Jason, 22  
Erk, Katrin, 27  
Erp, Marieke van, 28

- Fang, Hao, 20  
 Fang, Rui, 21  
 Feng, Yansong, 7, 16  
 Ferracane, Elisa, 27  
 Fidler, Sanja, 4  
 Finlayson, Mark, 28  
 Foreman, April, 25  
 Fraiberger, Samuel, 25  
 Fraser, Alex, 4  
 Freitas, Jesse, 24  
  
 Galley, Michel, 3  
 Ganchev, Kuzman, 4  
 Gogate, Vibhav, 7, 10  
 Goldberg, Yoav, 7, 8  
 Goldwasser, Dan, 22, 25  
 Goyal, Kartik, 22  
 Goyal, Naman, 22  
 Grau, Brigitte, 26  
 Grefenstette, Ed, 4  
 Grouin, Cyril, 26  
  
 Hamilton, William L., 25  
 Hamon, Thierry, 26  
 Handler, Abram, 25  
 Hathi, Shobhit, 21  
 Hermey, Yannik, 28  
 Hiemstra, Djoerd, 25  
 Hockenmaier, Julia, 4  
 Hollingshead, Kristy, 25  
 Hovy, Dirk, 24  
 Hovy, Dirk, 5  
 Hsieh, Yu-Lun, 21  
 Hsu, Wen-Lian, 21  
 Huang, Chung-Chi, 27  
 Huang, Degen, 22  
 Huang, Liang, 4  
 Huang, Ruihong, 4  
 Hwa, Rebecca, 3  
  
 Ifrim, Georgiana, 28  
 Ingraham, Loring, 25  
 Ireland, Molly, 25  
  
 Jaech, Aaron, 21  
 Jain, Siddharth, 20  
 Ji, Heng, 24  
 Johnson, Kristen, 25  
 Johnson, Todd, 26  
  
 Joseph, Kenneth, 24  
 Jurgens, David, 24  
  
 Kan, Min-Yen, 4  
 Kawahara, Daisuke, 4  
 Kim, Sunghwan Mac, 20, 25  
 Kirchhoff, Katrin, 26  
 Knight, Kevin, 23, 24  
 Knowles, Rebecca, 25  
 Kokkinakis, Dimitrios, 26  
 Kondrak, Greg, 3  
 Kriesi, Hanspeter, 25  
 Krishnan, Vinodh, 28  
 Ku, Lun-Wei, 20  
  
 Lalor, John, 26  
 Laverigne, Thomas, 27  
 Leskovec, Jure, 25  
 Li, Chang, 22  
 Li, Cheng-Te, 20  
 Li, Hang, 3  
 Li, Quanzhi, 21  
 Liu, Bing, 7, 12  
 LIU, XIAOMO, 21  
 Liu, Bing, 5  
 Liu, Yang, 4  
 Lorenzini, Jasmine, 25  
 Louis, Annie, 3  
 Loveys, Kate, 25  
 Lu, Zhiyong, 27  
 Lundholm Fors, Kristina, 26  
  
 Makarov, Peter, 25  
 Marcu, Daniel, 23  
 Marshall, Iain, 27  
 Martins, Andre, 4  
 McKeown, Kathy, 24  
 Mihalcea, Rada, 25  
 Miller, Ben, 28  
 Mohammad, Saif M., 3  
 Mohammad, Saif, 5  
 Montana, Giovanni, 26  
 Mooney, Ray, 4  
 Mukherjee, Arjun, 26  
 Mulcaire, George, 21  
 Munkhdalai, Tsendsuren, 26  
 Muresan, Smaranda, 5  
  
 Névél, Aurélie, 26

- Nakov, Preslav, 5  
 Nastase, Vivi, 5  
 Navindgi, Amit, 20  
 Nay, John J., 24  
 Need, Ariana, 25  
 Neubig, Graham, 7, 8  
 Neveol, Aurelie, 26  
 Ng, Vincent, 3, 7, 10  
 Nordlund, Arto, 26  
 Nourbakhsh, Armineh, 21  
 Nowson, Scott, 20  
  
 O'Connor, Brendan, 24, 25  
 O'Gorman, Tim, 28  
 Oh, Alice, 24  
 Orth, Carsten, 28  
 Ostendorf, Mari, 20, 21  
  
 Pacheco, María Leonor, 22  
 Palen, Leysia, 20  
 Palmer, Martha, 20, 28  
 Pappas, Nikolaos, 21  
 Paris, Cecile, 20, 25  
 Patwardhan, Siddharth, 3  
 Paul, Michael J., 20  
 Pedersen, Ted, 20, 26  
 Pennebaker, James, 25  
 Petersen, Steffen, 26  
 Piccardi, Massimo, 26  
 Pighin, Daniele, 3  
 Poghosyan, Gevorg, 28  
 Popescu-Belis, Andrei, 21  
 Poulston, Adam, 24  
 Priante, Anna, 25  
  
 Quattoni, Ariadna, 4  
  
 Rasmussen, Priscilla, 3  
 Reddy, Sravana, 24  
 Resnik, Rebecca, 25  
 Rey-Villamizar, Nicolas, 20, 26  
 Rheault, Ludovic, 25  
 Riedel, Sebastian, 20, 26  
 Rimell, Laura, 4  
 Ringger, Eric, 4  
 Ritter, Alan, 4  
 Roark, Brian, 5  
 Robert, Aude, 26  
 Robinson, Bella, 25  
  
 Rocktäschel, Tim, 20  
 Rosenthal, Sara, 24  
 Rush, Alexander, 22  
  
 Sadeque, Farig, 20, 26  
 Saeed, Aaqib, 25  
 Schneider, Darius, 28  
 Schwartz, H. Andrew, 25  
 Selzer, Stefan, 28  
 Shah, Sameena, 21  
 Shrestha, Prasha, 20, 26  
 Simonson, Dan, 28  
 Smith, Noah A., 21  
 Smith, David, 5  
 Solorio, Tamar, 20, 26  
 Somasundaran, Swapna, 3  
 Spanakis, Gerasimos, 28  
 Spithourakis, Georgios, 26  
 Srikumar, Vivek, 22  
 Stede, Manfred, 5  
 Stevenson, Mark, 24  
 Stevenson, Suzanne, 4  
 Stowe, Kevin, 20  
 Strube, Michael, 5  
 Su, Jian, 1–3  
 Sun, Xu, 7, 16  
  
 Tetreault, Joel, 5  
 Tin Vo, Duy, 13  
 Tran, Ke M., 23  
 Tsur, Oren, 24  
 Turner, Anne M., 26  
  
 Van Durme, Benjamin, 22  
 Vanderwende, Lucy, 5  
 Vaswani, Ashish, 23  
 Venugopal, Deepak, 7, 10  
 Vo, Duy Tin, 7, 13  
 Volkova, Svitlana, 24  
 Vossen, Piek, 28  
  
 Wallace, Byron, 26  
 Wallace, Byron C., 27  
 Wallach, Hanna, 25  
 Wan, Stephen, 20, 25  
 Wang, Alex, 25  
 Waseem, Zeerak, 25  
 Wen-Tau Yih, Scott, 4  
 Wilson, Steven, 25

Withey, Samuel, 26  
Wolfe, Travis, 22  
Wright-Bettner, Kristin, 28  
Wu, Dekai, 4  
  
Xu, Wei, 5  
  
Yang, Bishan, 3  
Yu, Hong, 26  
Yu, Zhiguo, 26  
Yung-jen Hsu, Jane, 20  
  
Zare Borzeshi, Ehsan, 26  
Zhang, Xiao, 22  
Zhang, Yue, 7, 13  
Zweig, Geoff, 5  
Zweigenbaum, Pierre, 26, 27